

# Task3

November 9, 2021

## 1 Notebook for task 3

We will start with some basic imports

```
[1]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
from scipy.stats import chi2
```

### 1.1 Part (a)

in order for the `bash` commands to execute, the paths for `plink` and `BOLT-LMM` have to be set accordingly (see the respective commands).

#### 1.1.1 Use `plink` to run a Cochran-Armitage test

```
[2]: !.../..plink/plink --bfile ../data/plink --model trend-only --out 3a --adjust
```

```
PLINK v1.90b6.24 64-bit (6 Jun 2021)          www.cog-genomics.org/plink/1.9/
(C) 2005-2021 Shaun Purcell, Christopher Chang GNU General Public License v3
Logging to 3a.log.
Options in effect:
  --adjust
  --bfile ../data/plink
  --model trend-only
  --out 3a

7905 MB RAM detected; reserving 3952 MB for main workspace.
1440616 variants loaded from .bim file.
619 people (305 males, 314 females) loaded from .fam.
619 phenotype values loaded from .fam.
Using 1 thread (no multithreaded calculations invoked).
Before main variant filters, 523 founders and 96 nonfounders present.
Calculating allele frequencies... 1011121314151617181920212223242526272829303132
33343536373839404142434445464748495051525354555657585960616263646566676869707172
737475767778798081828384858687888990919293949596979899 done.
Warning: 1879 het. haploid genotypes present (see 3a.hh ); many commands treat
these as missing.
```

Total genotyping rate is 0.99772.  
 1440616 variants and 619 people pass filters and QC.  
 Among remaining phenotypes, 119 are cases and 500 are controls.  
 Excluding 112 MT/haploid variants from --model analysis.  
 Writing --model report to 3a.model ... 10111213141516171819202122232425262728293  
 03132333435363738394041424344454647484950515253545556575859606162636465666768697  
 07172737475767778798081828384858687888990919293949596979899done.  
 --adjust: Genomic inflation est. lambda (based on median chisq) = 3.99938.  
 10111213141516171819202122232425262728293031323334353637383940414243444546474849  
 50515253545556575859606162636465666768697071727374757677787980818283848586878889  
 90919293949596979899--adjust values (1440504 variants) written to  
 3a.model.trend.adjusted .

### 1.1.2 Translate the space-separated output to tab-separated output which pandas can parse

```
[3]: !cat 3a.model | tr -s " " "\t" > 3a.tsv
```

### 1.1.3 Load into pandas and show the table

```
[4]: table = pd.read_csv("3a.tsv", sep="\t")
```

```
[5]: table.head()
```

```
[5]:
```

	Unnamed: 0	CHR	SNP	A1	A2	TEST	AFF	UNAFF	CHISQ	DF	\
0	NaN	1	rs10458597	T	C	TREND	4/230	13/973	0.1565	1.0	
1	NaN	1	rs2185539	T	C	TREND	7/231	0/1000	29.7500	1.0	
2	NaN	1	rs11240767	T	C	TREND	10/228	15/985	6.2060	1.0	
3	NaN	1	rs12564807	G	A	TREND	0/238	0/1000	NaN	NaN	
4	NaN	1	rs3131972	A	G	TREND	82/156	259/741	6.6180	1.0	

```

P
0 6.924000e-01
1 4.920000e-08
2 1.273000e-02
3 NaN
4 1.010000e-02


```

Here we can easily drop the *nan*-values


```
[6]: ps = table["P"].dropna().to_numpy()
```

Also we do some plots to get an idea of the data

```
[7]: plt.plot(sorted(ps))
plt.show()
```



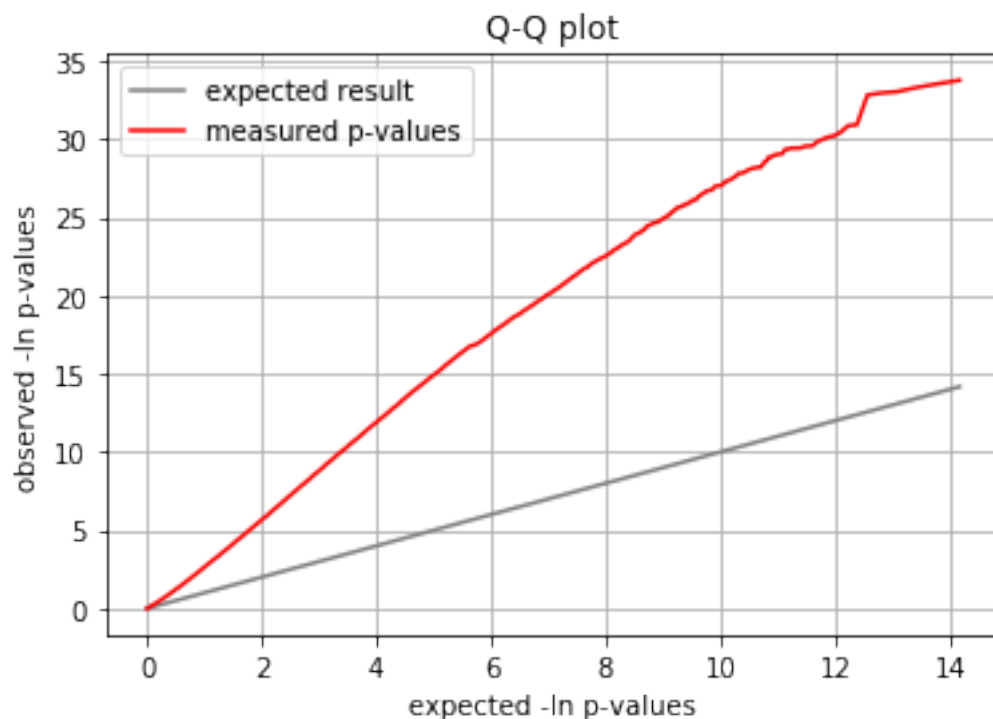
```
[8]: plt.hist(ps)  
plt.show()
```



## 1.2 Q-Q plot

First we generate the expected p-values (uniform distributed => `np.linspace`), then do the required `ln`-transform and plot

```
[9]: X = np.linspace(0,1,ps.size+1)[1:] # avoid to get a 0 since ln(0) is a problem
nln_X = -np.log(X)
nln_Y = -np.log(np.sort(ps))
plt.plot(nln_X, nln_X, color="grey", label="expected result")
plt.plot(nln_X, nln_Y, color="red", label="measured p-values")
plt.legend(loc="upper left")
plt.grid()
plt.xlabel("expected -ln p-values")
plt.ylabel("observed -ln p-values")
plt.title("Q-Q plot")
plt.show()
```



A Q-Q plot shows the measured p-values plotted against the expected ones. This should mostly be straight line like ( $y = x$ ) with possibly some exception towards the very small p-values. In order to make it easier so see that area one uses a negative `ln` transform. This makes the small values more pronounced and also puts them in the top right corner. As we can see from the plot already this is not calibrated at all, since the red and grey lines do not look anything alike.

This can also be seen from the `lambda`, which has been calculated by the first `plink` command: `> -adjust: Genomic inflation est. lambda (based on median chisq) = 3.99938.`

which is significantly larger than 1

### 1.3 Part (b)

We do not have to change the command from a since we have already been performing the adjusted test. We have, however, to also use the adjusted output:

```
[10]: !cat 3a.model.trend.adjusted | tr -s " " "\t" > 3a.adjusted.tsv
```

```
[11]: table_b = pd.read_csv("3a.adjusted.tsv", sep="\t")
```

```
[12]: table_b.head()
```

```
[12]:
```

	Unnamed: 0	CHR	SNP	UNADJ	GC	\
NaN	6	rs11962226	2.168000e-15	0.000073	3.123000e-09	
NaN	9	rs7031414	3.315000e-15	0.000082	4.775000e-09	
NaN	3	rs9311319	4.536000e-15	0.000089	6.534000e-09	
NaN	4	rs10007859	4.953000e-15	0.000091	7.135000e-09	
NaN	1	rs17130151	5.589000e-15	0.000094	8.050000e-09	

	BONF	HOLM	SIDAK_SS	SIDAK_SD	FDR_BH	\
NaN	3.123000e-09	3.123000e-09	3.123000e-09	1.610000e-09	2.376000e-08	
NaN	4.775000e-09	4.775000e-09	4.775000e-09	1.610000e-09	2.376000e-08	
NaN	6.534000e-09	6.534000e-09	6.534000e-09	1.610000e-09	2.376000e-08	
NaN	7.135000e-09	7.135000e-09	7.135000e-09	1.610000e-09	2.376000e-08	
NaN	8.050000e-09	8.050000e-09	8.050000e-09	1.610000e-09	2.376000e-08	

	FDR_BY
NaN	NaN
NaN	NaN
NaN	NaN
NaN	NaN
NaN	NaN


Here one can see the *GC* (genomic control) column, which has the adjusted values.

```
[13]: ps_gc = table_b["GC"].dropna().to_numpy()
      ps_gc[:5]
```

```
[13]: array([3.123e-09, 4.775e-09, 6.534e-09, 7.135e-09, 8.050e-09])
```

```
[14]: X = np.linspace(0,1,ps_gc.size+1)[1:] # avoid to get a 0 since ln(0) is a
      ↪problem
      nln_X = -np.log(X)
      nln_Y = -np.log(np.sort(ps_gc))
      plt.plot(nln_X, nln_X, color="grey", label="expected result")
      plt.plot(nln_X, nln_Y, color="red", label="measured p-values")
      plt.legend(loc="upper left")
```

```
plt.grid()
plt.xlabel("expected -ln p-values")
plt.ylabel("observed -ln p-values")
plt.title("Q-Q plot")
plt.show()
```



The plot still does not look good, but one can see, that the red line is now roughly centered with the grey one (it is above the grey one about as much as below). Still the red line does not follow the grey line as nicely as in the example in the slides.

#### 1.4 Part (c)

here we have to provide the *gc* value to the *-adjust* command. This causes *plink* to use the *gc*-adjusted values in the formulas. The result we are then looking for is in the *BONF* column.

```
[15]: !../plink/plink --bfile ../data/plink --model trend-only --out 3c --adjust gc
```

```
PLINK v1.90b6.24 64-bit (6 Jun 2021)          www.cog-genomics.org/plink/1.9/
(C) 2005-2021 Shaun Purcell, Christopher Chang GNU General Public License v3
Logging to 3c.log.
Options in effect:
  --adjust gc
  --bfile ../data/plink
  --model trend-only
```

```
--out 3c
```

```
7905 MB RAM detected; reserving 3952 MB for main workspace.
1440616 variants loaded from .bim file.
619 people (305 males, 314 females) loaded from .fam.
619 phenotype values loaded from .fam.
Using 1 thread (no multithreaded calculations invoked).
Before main variant filters, 523 founders and 96 nonfounders present.
Calculating allele frequencies... 1011121314151617181920212223242526272829303132
33343536373839404142434445464748495051525354555657585960616263646566676869707172
73747576777879808182838485868788899091929394959697989 done.
Warning: 1879 het. haploid genotypes present (see 3c.hh ); many commands treat
these as missing.
Total genotyping rate is 0.99772.
1440616 variants and 619 people pass filters and QC.
Among remaining phenotypes, 119 are cases and 500 are controls.
Excluding 112 MT/haploid variants from --model analysis.
Writing --model report to 3c.model ... 10111213141516171819202122232425262728293
03132333435363738394041424344454647484950515253545556575859606162636465666768697
07172737475767778798081828384858687888990919293949596979899done.
--adjust: Genomic inflation est. lambda (based on median chisq) = 3.99938.
10111213141516171819202122232425262728293031323334353637383940414243444546474849
50515253545556575859606162636465666768697071727374757677787980818283848586878889
90919293949596979899--adjust values (1440504 variants) written to
3c.model.trend.adjusted .
```

```
[16]: !cat 3c.model.trend.adjusted | tr -s " " "\t" > 3c.adjusted.tsv
```

```
[17]: table_c = pd.read_csv("3c.adjusted.tsv", sep="\t")
```

```
[18]: table_c.head()
```

```
[18]:
```


	Unnamed: 0	CHR	SNP	UNADJ	GC	BONF	HOLM	SIDAK_SS	\
NaN	6	rs11962226	2.168000e-15	0.000073	1	1	1	1	
NaN	9	rs7031414	3.315000e-15	0.000082	1	1	1	1	
NaN	3	rs9311319	4.536000e-15	0.000089	1	1	1	1	
NaN	4	rs10007859	4.953000e-15	0.000091	1	1	1	1	
NaN	1	rs17130151	5.589000e-15	0.000094	1	1	1	1	

	SIDAK_SD	FDR_BH	FDR_BY
NaN	0.9957	1	NaN
NaN	0.9957	1	NaN
NaN	0.9957	1	NaN
NaN	0.9957	1	NaN
NaN	0.9957	1	NaN

```
[19]: ps_gc_bonf = table_c["GC"].dropna().to_numpy()
```

```
[20]: X = np.linspace(0,1,ps_gc_bonf.size+1)[1:] # avoid to get a 0 since ln(0) is a
      ↪problem
      nln_X = -np.log(X)
      nln_Y = -np.log(np.sort(ps_gc_bonf))
      plt.plot(nln_X, nln_X, color="grey", label="expected result")
      plt.plot(nln_X, nln_Y, color="red", label="measured p-values")
      plt.legend(loc="upper left")
      plt.grid()
      plt.xlabel("expected -ln p-values")
      plt.ylabel("observed -ln p-values")
      plt.title("Q-Q plot")
      plt.show()
```



```
[21]: print(np.sort(ps_gc_bonf)[:10])
```

```
[1 1 1 1 1 1 1 1 1 1]
```

We can see, that all the p-values have been *corrected* up to 1. One possible explanation is, that we have very uncalibrated data ( $\lambda \approx 4$ ) and we do a lot of tests, so by the time both of those effects have been corrected, the results are just very uncertain (p-value of 1 is most uncertain).

### 1.5 Part (d)

This is probably not the proper way, to do it but it ended up giving reasonable results, so maybe it is not that far off after all.



The given data generates a problem, since there are chromosomes with numbers > 23 (and humans only have 23 chromosomes). So first we will note all the SNPs in chromosomes that are > 23. We cannot just delete them, since then there is a problem with incompatibilities (different lengths) between the .bim and .bed files

```
[22]: !cat ../data/plink.bim | grep "^2[3-9]" | cut -f2 > 3d.exclude
```

Since we cannot delete the bad SNPs we rename their chromosomes to ones that are actually valid (this should not cause problems, since we will exclude them anyways).

```
[23]: !cat ../data/plink.bim | sed "s/^2[3-9]/1/g" > plink.bim.23
```

Finally since there are too many SNPs for FastLMM to handle (it could do it, but there is a warning for more than 1 000 000 SNPs) we can prune some due to linkage disequilibrium (the idea is from the error message of FastLMM and the command from <https://www.cog-genomics.org/plink/1.9/ld>).

```
[24]: !../.. /plink/plink -bfile ../data/plink --indep-pairwise 50 5 0.5
```

```
PLINK v1.90b6.24 64-bit (6 Jun 2021)          www.cog-genomics.org/plink/1.9/
(C) 2005-2021 Shaun Purcell, Christopher Chang  GNU General Public License v3
Logging to plink.log.
Options in effect:
  --bfile ../data/plink
  --indep-pairwise 50 5 0.5
```

```
7905 MB RAM detected; reserving 3952 MB for main workspace.
1440616 variants loaded from .bim file.
619 people (305 males, 314 females) loaded from .fam.
619 phenotype values loaded from .fam.
Using 1 thread (no multithreaded calculations invoked).
Before main variant filters, 523 founders and 96 nonfounders present.
Calculating allele frequencies... 1011121314151617181920212223242526272829303132
33343536373839404142434445464748495051525354555657585960616263646566676869707172
73747576777879808182838485868788899091929394959697989 done.
Warning: 1879 het. haploid genotypes present (see plink.hh ); many commands
treat these as missing.
Total genotyping rate is 0.99772.
1440616 variants and 619 people pass filters and QC.
Among remaining phenotypes, 119 are cases and 500 are controls.
Pruned 75937 variants from chromosome 1, leaving 40628.
Pruned 77042 variants from chromosome 2, leaving 39540.
Pruned 63102 variants from chromosome 3, leaving 33569.
Pruned 55824 variants from chromosome 4, leaving 30076.
Pruned 57509 variants from chromosome 5, leaving 30556.
Pruned 60339 variants from chromosome 6, leaving 31161.
Pruned 48616 variants from chromosome 7, leaving 26819.
Pruned 49402 variants from chromosome 8, leaving 25975.
Pruned 40620 variants from chromosome 9, leaving 23086.
Pruned 48077 variants from chromosome 10, leaving 25859.
```

```

Pruned 46835 variants from chromosome 11, leaving 24257.
Pruned 43987 variants from chromosome 12, leaving 24652.
Pruned 33434 variants from chromosome 13, leaving 18577.
Pruned 28920 variants from chromosome 14, leaving 16637.
Pruned 26414 variants from chromosome 15, leaving 15998.
Pruned 27265 variants from chromosome 16, leaving 17445.
Pruned 23179 variants from chromosome 17, leaving 15278.
Pruned 25598 variants from chromosome 18, leaving 15293.
Pruned 15191 variants from chromosome 19, leaving 11072.
Pruned 22576 variants from chromosome 20, leaving 13726.
Pruned 11920 variants from chromosome 21, leaving 7411.
Pruned 11918 variants from chromosome 22, leaving 8191.
Pruned 34937 variants from chromosome 23, leaving 15572.
Pruned 138 variants from chromosome 25, leaving 346.
Pruned 47 variants from chromosome 26, leaving 65.
Pruning complete. 928827 of 1440616 variants removed.
Marker lists written to plink.prune.in and plink.prune.out .

```

Now with some SNPs pruned (we are now down to about 500 000), we can run the BOLT\_LMM command using the changes files from above (the --LDscoresUseChip again was given in an error message).

```

[25]: !../BOLT-LMM_v2.3.5/bolt --fam=../data/plink.fam --bim=plink.bim.23 --bed=../
      ↪data/plink.bed --phenoUseFam --lmm --statsFile 3d --LDscoresUseChip
      ↪--exclude 3d.exclude --exclude plink.prune.out --numThreads 8

```

```

+-----+
|      |      |      |
|  BOLT-LMM, v2.3.5  /_ /  |
|  March 20, 2021    /_ /  |
|  Po-Ru Loh         //   |
|                   /     |
+-----+

```

Copyright (C) 2014-2021 Harvard University.  
Distributed under the GNU GPLv3 open source license.

Compiled with USE\_SSE: fast aligned memory access  
Compiled with USE\_MKL: Intel Math Kernel Library linear algebra  
Boost version: 1\_58

Command line options:

```

../BOLT-LMM_v2.3.5/bolt \
--fam=../data/plink.fam \
--bim=plink.bim.23 \
--bed=../data/plink.bed \
--phenoUseFam \

```

```

--lmm \
--statsFile 3d \
--LDscoresUseChip \
--exclude 3d.exclude \
--exclude plink.prune.out \
--numThreads 8

Setting number of threads to 8
fam: ../data/plink.fam
bim(s): plink.bim.23
bed(s): ../data/plink.bed

=== Reading genotype data ===

Total indivs in PLINK data: Nbed = 619
Total indivs stored in memory: N = 619
Reading bim file #1: plink.bim.23
WARNING: Out-of-order snp in bim file: plink.bim.23
Line 1596:
1      AFX-SNP_11906976__rs581070      0      4289059 C      G
WARNING: Out-of-order snp in bim file: plink.bim.23
Line 2223:
1      AFX-SNP_6869948__rs10915322      0      4997214 A      G
WARNING: Out-of-order snp in bim file: plink.bim.23
Line 2273:
1      AFX-SNP_4600__rs9439505      0      5093533 T      C
WARNING: Out-of-order snp in bim file: plink.bim.23
Line 2320:
1      AFX-SNP_4992__rs7536712      0      5153524 T      C
WARNING: Out-of-order snp in bim file: plink.bim.23
Line 2445:
1      AFX-SNP_9025902__rs4847565      0      5284188 G      A
WARNING: Total number of out-of-order snps in bim file: 2051
      Read 1440616 snps
Total snps in PLINK data: Mbed = 1440616
Reading exclude file (SNPs to exclude): 3d.exclude
Excluded 51105 SNP(s)
Reading exclude file (SNPs to exclude): plink.prune.out
Excluded 893705 SNP(s)

Breakdown of SNP pre-filtering results:
  495806 SNPs to include in model (i.e., GRM)
  0 additional non-GRM SNPs loaded
  944810 excluded SNPs
WARNING: No genetic map provided; using physical positions only
Allocating 495806 x 620/4 bytes to store genotypes
Reading genotypes and performing QC filtering on snps and indivs...
Reading bed file #1: ../data/plink.bed

```

```

    Expecting 223295480 (+3) bytes for 619 indivs, 1440616 snps
Total indivs after QC: 619
Total post-QC SNPs: M = 495806
    Variance component 1: 495806 post-QC SNPs (name: 'modelSnps')
Time for SnpData setup = 10.2737 sec

=== Reading phenotype and covariate data ===

Number of indivs with no missing phenotype(s) to use: 619
NOTE: Using all-1s vector (constant term) in addition to specified covariates
    Using quantitative covariate: CONST_ALL_ONES
Number of individuals used in analysis: Nused = 619
Singular values of covariate matrix:
    S[0] = 24.8797
Total covariate vectors: C = 1
Total independent covariate vectors: Cind = 1

=== Initializing Bolt object: projecting and normalizing SNPs ===

Number of chroms with >= 1 good SNP: 22
Average norm of projected SNPs: 618.000000
Dimension of all-1s proj space (Nused-1): 618
Time for covariate data setup + Bolt initialization = 1.68237 sec

Phenotype 1:  N = 619  mean = 1.19225  std = 0.394384

=== Computing linear regression (LINREG) stats ===

Time for computing LINREG stats = 0.389594 sec

=== Estimating variance parameters ===

Using CGtol of 0.005 for this step
Using default number of random trials: 15 (for Nused = 619)

Estimating MC scaling f_REML at log(delta) = 1.09861, h2 = 0.25...
    Batch-solving 16 systems of equations using conjugate gradient iteration
    iter 1:  time=2.34  rNorms/orig: (0.6,2)  res2s: 1097.46..18.0129
    iter 2:  time=1.71  rNorms/orig: (0.1,0.5)  res2s: 1197.42..40.9796
    iter 3:  time=1.91  rNorms/orig: (0.03,0.2)  res2s: 1225.58..45.2024
    iter 4:  time=1.65  rNorms/orig: (0.02,0.04)  res2s: 1263.1..48.3579
    iter 5:  time=1.76  rNorms/orig: (0.003,0.008)  res2s: 1267.37..48.4496
    iter 6:  time=1.91  rNorms/orig: (0.0007,0.002)  res2s: 1267.71..48.4583
    Converged at iter 6: rNorms/orig all < CGtol=0.005
    Time breakdown: dgemm = 55.8%, memory/overhead = 44.2%
    MCscaling: logDelta = 1.10, h2 = 0.250, f = 0.203599

Estimating MC scaling f_REML at log(delta) = 0, h2 = 0.5...

```

Batch-solving 16 systems of equations using conjugate gradient iteration  
iter 1: time=1.73 rNorms/orig: (2,2) res2s: 152.875..3.95465  
iter 2: time=1.62 rNorms/orig: (0.3,1) res2s: 242.219..15.0849  
iter 3: time=1.73 rNorms/orig: (0.2,0.4) res2s: 259.345..18.4906  
iter 4: time=1.81 rNorms/orig: (0.07,0.2) res2s: 299.292..23.5421  
iter 5: time=1.81 rNorms/orig: (0.03,0.05) res2s: 307.252..23.948  
iter 6: time=1.76 rNorms/orig: (0.01,0.02) res2s: 309.136..24.0482  
iter 7: time=1.70 rNorms/orig: (0.004,0.007) res2s: 309.566..24.069  
iter 8: time=1.65 rNorms/orig: (0.001,0.002) res2s: 309.623..24.0724  
Converged at iter 8: rNorms/orig all < CGtol=0.005  
Time breakdown: dgemm = 55.0%, memory/overhead = 45.0%  
MCscaling: logDelta = 0.00, h2 = 0.500, f = 0.0719073

Estimating MC scaling f\_REML at log(delta) = -0.599873, h2 = 0.645627...

Batch-solving 16 systems of equations using conjugate gradient iteration  
iter 1: time=1.73 rNorms/orig: (2,2) res2s: 39.9569..1.4435  
iter 2: time=1.85 rNorms/orig: (0.4,1) res2s: 85.6996..6.84329  
iter 3: time=1.64 rNorms/orig: (0.3,0.6) res2s: 95.6316..8.868  
iter 4: time=1.69 rNorms/orig: (0.1,0.2) res2s: 120.508..12.7935  
iter 5: time=1.79 rNorms/orig: (0.05,0.09) res2s: 126.586..13.2451  
iter 6: time=1.69 rNorms/orig: (0.03,0.04) res2s: 128.556..13.3987  
iter 7: time=1.82 rNorms/orig: (0.01,0.02) res2s: 129.301..13.4468  
iter 8: time=1.64 rNorms/orig: (0.004,0.006) res2s: 129.478..13.4596  
iter 9: time=1.73 rNorms/orig: (0.002,0.002) res2s: 129.509..13.4629  
Converged at iter 9: rNorms/orig all < CGtol=0.005  
Time breakdown: dgemm = 57.0%, memory/overhead = 43.0%  
MCscaling: logDelta = -0.60, h2 = 0.646, f = 0.0268059

Estimating MC scaling f\_REML at log(delta) = -0.956406, h2 = 0.722402...

Batch-solving 16 systems of equations using conjugate gradient iteration  
iter 1: time=1.74 rNorms/orig: (2,3) res2s: 17.4538..0.762534  
iter 2: time=2.01 rNorms/orig: (0.4,1) res2s: 44.3168..3.98776  
iter 3: time=4.47 rNorms/orig: (0.4,0.8) res2s: 50.7492..5.31301  
iter 4: time=2.55 rNorms/orig: (0.1,0.2) res2s: 67.2869..8.22528  
iter 5: time=1.82 rNorms/orig: (0.07,0.1) res2s: 71.7183..8.61819  
iter 6: time=1.87 rNorms/orig: (0.04,0.05) res2s: 73.2865..8.77167  
iter 7: time=1.71 rNorms/orig: (0.02,0.03) res2s: 74.0444..8.82938  
iter 8: time=1.98 rNorms/orig: (0.008,0.01) res2s: 74.2836..8.84875  
iter 9: time=2.27 rNorms/orig: (0.003,0.005) res2s: 74.3373..8.85495  
Converged at iter 9: rNorms/orig all < CGtol=0.005  
Time breakdown: dgemm = 49.3%, memory/overhead = 50.7%  
MCscaling: logDelta = -0.96, h2 = 0.722, f = 0.00817042

Estimating MC scaling f\_REML at log(delta) = -1.11272, h2 = 0.752636...

Batch-solving 16 systems of equations using conjugate gradient iteration  
iter 1: time=1.46 rNorms/orig: (2,3) res2s: 12.1377..0.57247  
iter 2: time=1.50 rNorms/orig: (0.5,2) res2s: 32.9609..3.10373  
iter 3: time=1.49 rNorms/orig: (0.4,0.8) res2s: 38.1615..4.17945

```

iter 4:  time=1.58  rNorms/orig: (0.2,0.3)  res2s: 51.6553..6.65785
iter 5:  time=1.49  rNorms/orig: (0.07,0.1)  res2s: 55.4047..7.01248
iter 6:  time=1.89  rNorms/orig: (0.04,0.06)  res2s: 56.7528..7.15814
iter 7:  time=1.50  rNorms/orig: (0.02,0.04)  res2s: 57.4706..7.21677
iter 8:  time=1.46  rNorms/orig: (0.01,0.01)  res2s: 57.7236..7.23824
iter 9:  time=1.54  rNorms/orig: (0.004,0.006)  res2s: 57.7858..7.24574
iter 10: time=1.51  rNorms/orig: (0.002,0.002)  res2s: 57.7985..7.24713
Converged at iter 10: rNorms/orig all < CGtol=0.005
Time breakdown: dgemm = 56.3%, memory/overhead = 43.7%
MCscaling: logDelta = -1.11, h2 = 0.753, f = 0.00153066

```

```

Estimating MC scaling f_REML at log(delta) = -1.14876, h2 = 0.759284...
Batch-solving 16 systems of equations using conjugate gradient iteration
iter 1:  time=1.64  rNorms/orig: (2,3)  res2s: 11.165..0.535585
iter 2:  time=1.61  rNorms/orig: (0.5,2)  res2s: 30.7718..2.92634
iter 3:  time=1.36  rNorms/orig: (0.4,0.8)  res2s: 35.7152..3.94973
iter 4:  time=1.55  rNorms/orig: (0.2,0.3)  res2s: 48.5667..6.3319
iter 5:  time=1.46  rNorms/orig: (0.07,0.1)  res2s: 52.1667..6.6771
iter 6:  time=1.41  rNorms/orig: (0.04,0.06)  res2s: 53.4634..6.82043
iter 7:  time=1.45  rNorms/orig: (0.02,0.04)  res2s: 54.1688..6.87898
iter 8:  time=1.37  rNorms/orig: (0.01,0.01)  res2s: 54.4236..6.90083
iter 9:  time=1.47  rNorms/orig: (0.005,0.006)  res2s: 54.4874..6.90861
iter 10: time=1.50  rNorms/orig: (0.002,0.003)  res2s: 54.5007..6.91007
Converged at iter 10: rNorms/orig all < CGtol=0.005
Time breakdown: dgemm = 58.7%, memory/overhead = 41.3%
MCscaling: logDelta = -1.15, h2 = 0.759, f = 0.000109975

```

```

Secant iteration for h2 estimation converged in 4 steps
Estimated (pseudo-)heritability: h2g = 0.759
To more precisely estimate variance parameters and estimate s.e., use --reml
Variance params: sigma^2_K = 0.116254, logDelta = -1.148757, f = 0.000109975

```

Time for fitting variance components = 98.566 sec

=== Computing mixed model assoc stats (inf. model) ===

```

Selected 30 SNPs for computation of prospective stat
Tried 30; threw out 0 with GRAMMAR chisq > 5
Assigning SNPs to 22 chunks for leave-out analysis
Each chunk is excluded when testing SNPs belonging to the chunk
Batch-solving 52 systems of equations using conjugate gradient iteration
iter 1:  time=2.73  rNorms/orig: (1,3)  res2s: 0.538387..1.11321
iter 2:  time=2.98  rNorms/orig: (0.2,2)  res2s: 2.92479..25.5659
iter 3:  time=3.05  rNorms/orig: (0.2,0.9)  res2s: 3.94232..28.904
iter 4:  time=2.95  rNorms/orig: (0.02,0.3)  res2s: 6.31714..37.1859
iter 5:  time=3.03  rNorms/orig: (0.01,0.1)  res2s: 6.65944..40.2549
iter 6:  time=2.83  rNorms/orig: (0.006,0.06)  res2s: 6.8044..41.2866
iter 7:  time=3.11  rNorms/orig: (0.002,0.03)  res2s: 6.86191..41.6184

```

```

iter 8:  time=3.19  rNorms/orig: (0.0009,0.01)  res2s: 6.88329..41.755
iter 9:  time=3.08  rNorms/orig: (0.0005,0.006)  res2s: 6.89117..41.7828
iter 10: time=3.22  rNorms/orig: (0.0002,0.002)  res2s: 6.89261..41.7888
iter 11: time=3.03  rNorms/orig: (0.0001,0.0009)  res2s: 6.89291..41.79
iter 12: time=3.08  rNorms/orig: (0.0002,0.002)  res2s: 6.89296..41.7903
iter 13: time=2.91  rNorms/orig: (4e-05,0.0008)  res2s: 6.89296..41.7902
iter 14: time=2.82  rNorms/orig: (2e-05,0.0002)  res2s: 6.89297..41.7903
Converged at iter 14: rNorms/orig all < CGtol=0.0005
Time breakdown: dgemm = 69.9%, memory/overhead = 30.1%

AvgPro: 0.828  AvgRetro: 0.592  Calibration: 1.398 (0.096)  (30 SNPs)
Ratio of medians: 1.212  Median of ratios: 1.298
WARNING: Calibration std error is high; consider increasing --numCalibSnps
        Using ratio of medians instead: 1.21167

Time for computing infinitesimal model assoc stats = 42.764 sec

=== Estimating chip LD Scores using 400 indivs ===

Time for estimating chip LD Scores = 2.20954 sec

WARNING: No LDscoresFile provided; using estimated LD among chip SNPs

=== Estimating mixture parameters by cross-validation ===

Setting maximum number of iterations to 250 for this step
Max CV folds to compute = 5 (to have > 10000 samples)

====> Starting CV fold 1 <====

NOTE: Using all-1s vector (constant term) in addition to specified covariates
      Using quantitative covariate: CONST_ALL_ONES
Number of individuals used in analysis: Nused = 495
Singular values of covariate matrix:
      S[0] = 22.2486
Total covariate vectors: C = 1
Total independent covariate vectors: C indep = 1

=== Initializing Bolt object: projecting and normalizing SNPs ===

Number of chroms with >= 1 good SNP: 22
Average norm of projected SNPs: 494.000000
Dimension of all-1s proj space (Nused-1): 494
Beginning variational Bayes
iter 1:  time=5.55 for 18 active reps
iter 2:  time=4.19 for 18 active reps  approxLL diffs: (209.55,297.81)
iter 3:  time=3.61 for 18 active reps  approxLL diffs: (31.26,55.38)

```

```

iter 4:  time=3.94 for 18 active reps  approxLL diffs: (12.37,17.69)
iter 5:  time=3.89 for 18 active reps  approxLL diffs: (6.33,8.07)
iter 6:  time=3.66 for 18 active reps  approxLL diffs: (3.77,4.35)
iter 7:  time=3.84 for 18 active reps  approxLL diffs: (2.46,2.65)
iter 8:  time=3.90 for 18 active reps  approxLL diffs: (1.72,1.79)
iter 9:  time=3.68 for 18 active reps  approxLL diffs: (1.25,1.31)
iter 10: time=3.59 for 18 active reps  approxLL diffs: (0.91,0.99)
iter 11: time=3.55 for 18 active reps  approxLL diffs: (0.68,0.74)
iter 12: time=3.75 for 18 active reps  approxLL diffs: (0.53,0.57)
iter 13: time=3.92 for 18 active reps  approxLL diffs: (0.44,0.46)
iter 14: time=3.61 for 18 active reps  approxLL diffs: (0.37,0.39)
iter 15: time=3.55 for 18 active reps  approxLL diffs: (0.31,0.33)
iter 16: time=4.05 for 18 active reps  approxLL diffs: (0.27,0.28)
iter 17: time=4.14 for 18 active reps  approxLL diffs: (0.23,0.24)
iter 18: time=4.10 for 18 active reps  approxLL diffs: (0.20,0.21)
iter 19: time=3.74 for 18 active reps  approxLL diffs: (0.18,0.19)
iter 20: time=4.48 for 18 active reps  approxLL diffs: (0.16,0.17)
iter 21: time=3.83 for 18 active reps  approxLL diffs: (0.15,0.16)
iter 22: time=4.70 for 18 active reps  approxLL diffs: (0.13,0.15)
iter 23: time=3.76 for 18 active reps  approxLL diffs: (0.12,0.13)
iter 24: time=3.74 for 18 active reps  approxLL diffs: (0.11,0.12)
iter 25: time=3.74 for 18 active reps  approxLL diffs: (0.11,0.12)
iter 26: time=4.63 for 18 active reps  approxLL diffs: (0.10,0.11)
iter 27: time=3.98 for 18 active reps  approxLL diffs: (0.09,0.10)
iter 28: time=4.50 for 18 active reps  approxLL diffs: (0.09,0.10)
iter 29: time=5.13 for 18 active reps  approxLL diffs: (0.08,0.09)
iter 30: time=4.25 for 18 active reps  approxLL diffs: (0.08,0.08)
iter 31: time=3.87 for 18 active reps  approxLL diffs: (0.07,0.08)
iter 32: time=3.66 for 18 active reps  approxLL diffs: (0.07,0.08)
iter 33: time=3.61 for 18 active reps  approxLL diffs: (0.07,0.07)
iter 34: time=3.60 for 18 active reps  approxLL diffs: (0.06,0.07)
iter 35: time=3.65 for 18 active reps  approxLL diffs: (0.06,0.07)
iter 36: time=3.65 for 18 active reps  approxLL diffs: (0.06,0.06)
iter 37: time=3.50 for 18 active reps  approxLL diffs: (0.05,0.06)
iter 38: time=3.49 for 18 active reps  approxLL diffs: (0.05,0.06)
iter 39: time=3.39 for 18 active reps  approxLL diffs: (0.05,0.05)
iter 40: time=3.42 for 18 active reps  approxLL diffs: (0.05,0.05)
iter 41: time=3.52 for 18 active reps  approxLL diffs: (0.05,0.05)
iter 42: time=3.41 for 18 active reps  approxLL diffs: (0.04,0.05)
iter 43: time=3.34 for 18 active reps  approxLL diffs: (0.04,0.05)
iter 44: time=3.79 for 18 active reps  approxLL diffs: (0.04,0.04)
iter 45: time=3.57 for 18 active reps  approxLL diffs: (0.04,0.04)
iter 46: time=3.52 for 18 active reps  approxLL diffs: (0.04,0.04)
iter 47: time=3.72 for 18 active reps  approxLL diffs: (0.04,0.04)
iter 48: time=3.64 for 18 active reps  approxLL diffs: (0.04,0.04)
iter 49: time=3.70 for 18 active reps  approxLL diffs: (0.04,0.04)
iter 50: time=3.65 for 18 active reps  approxLL diffs: (0.03,0.04)
iter 51: time=3.59 for 18 active reps  approxLL diffs: (0.03,0.03)

```



```

iter 52:  time=3.69 for 18 active reps approxLL diffs: (0.03,0.03)
iter 53:  time=3.65 for 18 active reps approxLL diffs: (0.03,0.03)
iter 54:  time=3.59 for 18 active reps approxLL diffs: (0.03,0.03)
iter 55:  time=3.64 for 18 active reps approxLL diffs: (0.03,0.03)
iter 56:  time=3.67 for 18 active reps approxLL diffs: (0.03,0.03)
iter 57:  time=3.63 for 18 active reps approxLL diffs: (0.03,0.03)
iter 58:  time=3.62 for 18 active reps approxLL diffs: (0.03,0.03)
iter 59:  time=3.49 for 18 active reps approxLL diffs: (0.03,0.03)
iter 60:  time=3.65 for 18 active reps approxLL diffs: (0.02,0.03)
iter 61:  time=3.62 for 18 active reps approxLL diffs: (0.02,0.03)
iter 62:  time=3.68 for 18 active reps approxLL diffs: (0.02,0.03)
iter 63:  time=3.74 for 18 active reps approxLL diffs: (0.02,0.02)
iter 64:  time=3.71 for 18 active reps approxLL diffs: (0.02,0.02)
iter 65:  time=3.69 for 18 active reps approxLL diffs: (0.02,0.02)
iter 66:  time=3.66 for 18 active reps approxLL diffs: (0.02,0.02)
iter 67:  time=3.68 for 18 active reps approxLL diffs: (0.02,0.02)
iter 68:  time=3.61 for 18 active reps approxLL diffs: (0.02,0.02)
iter 69:  time=3.57 for 18 active reps approxLL diffs: (0.02,0.02)
iter 70:  time=3.62 for 18 active reps approxLL diffs: (0.02,0.02)
iter 71:  time=3.93 for 18 active reps approxLL diffs: (0.02,0.02)
iter 72:  time=3.87 for 18 active reps approxLL diffs: (0.02,0.02)
iter 73:  time=3.84 for 18 active reps approxLL diffs: (0.02,0.02)
iter 74:  time=3.61 for 18 active reps approxLL diffs: (0.02,0.02)
iter 75:  time=3.67 for 18 active reps approxLL diffs: (0.02,0.02)
iter 76:  time=3.72 for 18 active reps approxLL diffs: (0.02,0.02)
iter 77:  time=4.15 for 18 active reps approxLL diffs: (0.02,0.02)
iter 78:  time=3.57 for 18 active reps approxLL diffs: (0.02,0.02)
iter 79:  time=3.79 for 18 active reps approxLL diffs: (0.02,0.02)
iter 80:  time=3.52 for 18 active reps approxLL diffs: (0.02,0.02)
iter 81:  time=3.68 for 18 active reps approxLL diffs: (0.02,0.02)
iter 82:  time=3.77 for 18 active reps approxLL diffs: (0.02,0.02)
iter 83:  time=3.62 for 18 active reps approxLL diffs: (0.02,0.02)
iter 84:  time=3.60 for 18 active reps approxLL diffs: (0.02,0.02)
iter 85:  time=3.70 for 18 active reps approxLL diffs: (0.02,0.02)
iter 86:  time=3.82 for 18 active reps approxLL diffs: (0.01,0.02)
iter 87:  time=4.78 for 18 active reps approxLL diffs: (0.01,0.02)
iter 88:  time=4.59 for 18 active reps approxLL diffs: (0.01,0.01)
iter 89:  time=5.28 for 18 active reps approxLL diffs: (0.01,0.01)
iter 90:  time=4.28 for 18 active reps approxLL diffs: (0.01,0.01)
iter 91:  time=4.56 for 18 active reps approxLL diffs: (0.01,0.01)
iter 92:  time=3.98 for 18 active reps approxLL diffs: (0.01,0.01)
iter 93:  time=4.24 for 18 active reps approxLL diffs: (0.01,0.01)
iter 94:  time=3.40 for 18 active reps approxLL diffs: (0.01,0.01)
iter 95:  time=3.38 for 18 active reps approxLL diffs: (0.01,0.01)
iter 96:  time=3.44 for 18 active reps approxLL diffs: (0.01,0.01)
iter 97:  time=3.51 for 18 active reps approxLL diffs: (0.01,0.01)
iter 98:  time=3.44 for 18 active reps approxLL diffs: (0.01,0.01)
iter 99:  time=3.38 for 18 active reps approxLL diffs: (0.01,0.01)

```

```

iter 100: time=3.34 for 18 active reps approxLL diffs: (0.01,0.01)
iter 101: time=3.30 for 18 active reps approxLL diffs: (0.01,0.01)
iter 102: time=3.35 for 18 active reps approxLL diffs: (0.01,0.01)
iter 103: time=3.36 for 18 active reps approxLL diffs: (0.01,0.01)
iter 104: time=3.38 for 18 active reps approxLL diffs: (0.01,0.01)
iter 105: time=3.43 for 18 active reps approxLL diffs: (0.01,0.01)
iter 106: time=3.44 for 18 active reps approxLL diffs: (0.01,0.01)
iter 107: time=3.29 for 18 active reps approxLL diffs: (0.01,0.01)
iter 108: time=3.56 for 18 active reps approxLL diffs: (0.01,0.01)
iter 109: time=3.74 for 18 active reps approxLL diffs: (0.01,0.01)
iter 110: time=3.48 for 18 active reps approxLL diffs: (0.01,0.01)
iter 111: time=3.58 for 18 active reps approxLL diffs: (0.01,0.01)
iter 112: time=3.23 for 17 active reps approxLL diffs: (0.01,0.01)
iter 113: time=0.67 for 1 active reps approxLL diffs: (0.01,0.01)
iter 114: time=0.69 for 1 active reps approxLL diffs: (0.01,0.01)
Converged at iter 114: approxLL diffs each have been < LLtol=0.01
Time breakdown: dgemm = 32.7%, memory/overhead = 67.3%
Computing predictions on left-out cross-validation fold
Time for computing predictions = 1.80839 sec

```

Average PVEs obtained by param pairs tested (high to low):

```

f2=0.5, p=0.5: 0.301579
f2=0.3, p=0.5: 0.301555
f2=0.5, p=0.2: 0.301496
...
f2=0.1, p=0.01: 0.296609

```

Detailed CV fold results:

```

Absolute prediction MSE baseline (covariates only): 0.146285
Absolute prediction MSE using standard LMM:         0.102169
Absolute prediction MSE, fold-best f2=0.5, p=0.5: 0.102169
  Absolute pred MSE using f2=0.5, p=0.5: 0.102169
  Absolute pred MSE using f2=0.5, p=0.2: 0.102181
  Absolute pred MSE using f2=0.5, p=0.1: 0.102207
  Absolute pred MSE using f2=0.5, p=0.05: 0.102259
  Absolute pred MSE using f2=0.5, p=0.02: 0.102405
  Absolute pred MSE using f2=0.5, p=0.01: 0.102604
  Absolute pred MSE using f2=0.3, p=0.5: 0.102172
  Absolute pred MSE using f2=0.3, p=0.2: 0.102203
  Absolute pred MSE using f2=0.3, p=0.1: 0.102254
  Absolute pred MSE using f2=0.3, p=0.05: 0.102355
  Absolute pred MSE using f2=0.3, p=0.02: 0.102610
  Absolute pred MSE using f2=0.3, p=0.01: 0.102860
  Absolute pred MSE using f2=0.1, p=0.5: 0.102183
  Absolute pred MSE using f2=0.1, p=0.2: 0.102235
  Absolute pred MSE using f2=0.1, p=0.1: 0.102320
  Absolute pred MSE using f2=0.1, p=0.05: 0.102479
  Absolute pred MSE using f2=0.1, p=0.02: 0.102806

```

```

Absolute pred MSE using f2=0.1, p=0.01: 0.102896

====> End CV fold 1: 18 remaining param pair(s) <====

Estimated proportion of variance explained using inf model: 0.302
Relative improvement in prediction MSE using non-inf model: 0.000

Exiting CV: non-inf model does not substantially improve prediction
Optimal mixture parameters according to CV: f2 = 0.5, p = 0.5
Bayesian non-infinitesimal model does not fit substantially better
=> Not computing non-inf assoc stats (to override, use --lmmForceNonInf)

Time for estimating mixture parameters = 426.859 sec

Calibration stats: mean and lambdaGC (over SNPs used in GRM)
(note that both should be >1 because of polygenicity)
Mean BOLT_LMM_INF: 0.887406 (495806 good SNPs) lambdaGC: 0.806821

=== Streaming genotypes to compute and write assoc stats at all SNPs ===

Time for streaming genotypes and writing output = 8.53651 sec

Total elapsed time for analysis = 591.281 sec

```

### 1.5.1 Analysis of the result in python

Now we can import the generated file into pandas and have a look at it

```
[26]: lmm_data = pd.read_csv("3d", sep="\t")
```

```
[27]: lmm_data.head()
```

```
[27]:
```

	SNP	CHR	BP	GENPOS	ALLELE1	ALLELE0	A1FREQ	F_MISS	\
0	rs10458597	1	554484	0	T	C	0.013934	0.014540	
1	rs2185539	1	556738	0	T	C	0.005654	0.000000	
2	rs11240767	1	718814	0	T	C	0.020194	0.000000	
3	rs3131969	1	744045	0	A	G	0.277414	0.012924	
4	rs12562034	1	758311	0	A	G	0.258592	0.012924	

	BETA	SE	P_BOLT_LMM_INF
0	-0.006496	0.080937	0.94
1	0.194531	0.146021	0.18
2	0.004044	0.072609	0.96
3	-0.000252	0.022920	0.99
4	-0.022805	0.023099	0.32

For a quick check we can have a look at the smallest (strongest) p-values, which are pretty small so there is a chance this is good.

```
[28]: lmm_data.sort_values("P_BOLT_LMM_INF").head()
```

```
[28]:
```

	SNP	CHR	BP	GENPOS	ALLELE1	ALLELE0	A1FREQ	\
130871	rs17042171	4	111927736	0	A	C	0.297254	
402721	rs17158372	15	81144053	0	G	A	0.019418	
410483	rs12444503	16	10696526	0	A	G	0.004854	
487969	rs8190314	22	16606771	0	A	G	0.012195	
291319	rs4253212	10	50348218	0	A	G	0.039580	

	F_MISS	BETA	SE	P_BOLT_LMM_INF
130871	0.000000	0.140726	0.022166	2.200000e-10
402721	0.001616	0.453361	0.076621	3.300000e-09
410483	0.001616	0.839470	0.157517	9.900000e-08
487969	0.006462	0.528830	0.100227	1.300000e-07
291319	0.000000	0.292594	0.055867	1.600000e-07

Since we already have all the data in Python, we just quickly do the Bonferroni correction here (i.e. multiply the p-values by the number of tests).

```
[29]: n = lmm_data.size
print(n)
```

```
5453866
```

```
[30]: p_corrected = np.clip(lmm_data.P_BOLT_LMM_INF.to_numpy() * n, 0, 1)
```

```
[31]: for i, p in enumerate(np.sort(p_corrected)[:10]):
        print(f"The {i+1}-smallest p-value is {(p*100):2.1f}% and therefore {'' if_
        ↪p < 0.05 else 'not '}significant")
```

```
The 1-smallest p-value is 0.1% and therefore significant
The 2-smallest p-value is 1.8% and therefore significant
The 3-smallest p-value is 54.0% and therefore not significant
The 4-smallest p-value is 70.9% and therefore not significant
The 5-smallest p-value is 87.3% and therefore not significant
The 6-smallest p-value is 100.0% and therefore not significant
The 7-smallest p-value is 100.0% and therefore not significant
The 8-smallest p-value is 100.0% and therefore not significant
The 9-smallest p-value is 100.0% and therefore not significant
The 10-smallest p-value is 100.0% and therefore not significant
```

These are the smallest p-values after correction. We can see, that the first two are smaller than 5%, so we have found two significant results.

### 1.5.2 Comparison

These p-values are much better than the ones from before. Before we had (after calibration and correction for multiple testing) all p-values 1. This is obviously useless. Before these corrections we had so many very small p-values, that it is unrealistic, that this is correct.

Now we get two significant p-values, which are not overly small either, and everything else is much bigger and therefore insignificant. This appears much more reasonable.